

Building multi-region & multi-cloud services with Kafka

Travis Jeffery

- github.com/travisjeffery
- twitter.com/travisjeffery
- medium.com/@travisjeffery

What we're talking about today

- Building services on multiple regions, multiple clouds in a manageable way
- Example: running services in AWS with others in GCP or us-west-2 & us-east-1
- 3 system architectures according to message ingress and egress

Why?

Why multiple regions?

- **Latency:** Customers want your service to be fast.
- **Compliance:** You want a global customer base but different countries have different laws saying where/how data is stored.
- **Availability:** Increase your resilience above what a single region's availability zones.


Why multiple clouds?

- Features and services available differ on each cloud.
- Save money with better deals on different vendors.
- Reduce dependency on your provider.


How?

**A typical,
naive system**

Request-response to each region & cloud



Request-response to each region & cloud



Request-response to each region & cloud

Bad:


- Security: Secure all endpoints
- Routing: VPN tunnels, firewalls
- Service discovery: Must know where & how to msg

Good:

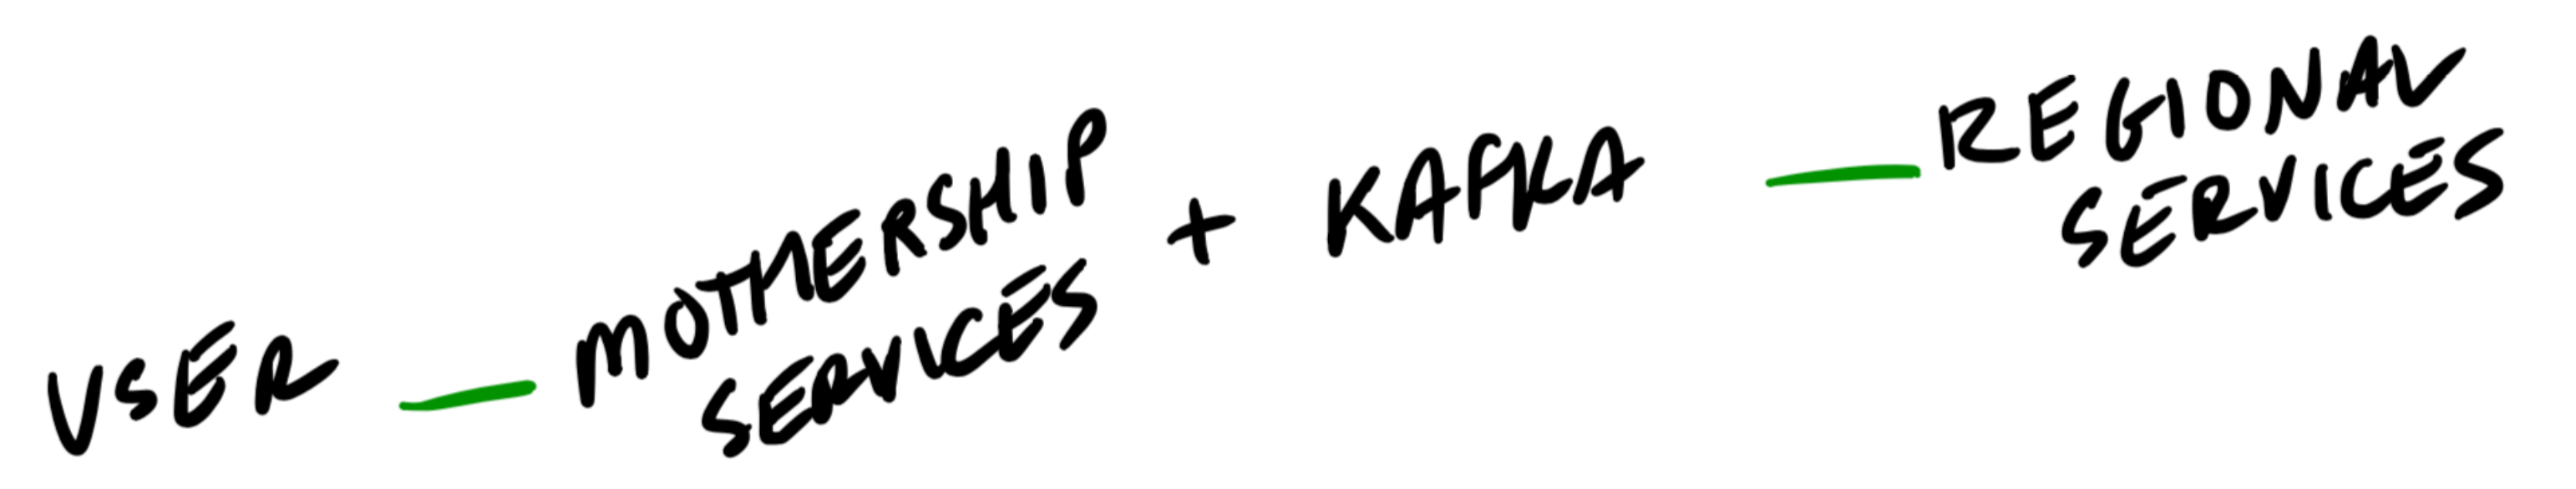
- Latency: Direct connections.

**A new,
smarter
system**

Using Kafka as central message hub



Message header routing: low ingress, low egress



Message header routing: low ingress, low egress

```
MSG = {  
  ROUTER = {  
    CLOUD: 'AWS',  
    REGION: 'US-WEST-2'  
  }  
  DATA: ...  
}
```

→ CREATE_ACCOUNT_TOPIC

```
REGIONAL SERVICE  
MSG ← CONSUMER  
IF MSG IS FOR ME  
  OPERATE  
ELSE  
  IGNORE
```

Topic per cloud and region: high ingress, low egress

MSG = {
 DATA: ...
}

→ CREATE_ACCOUNT_AWS_US_WEST_2

Topic per cloud and region into local kafka clusters: high ingress, high egress



Extending for when you need a request-response

- Problem: End-to-end latency, for example a UI with a live search field
- Solution: Give the UI a regional API to request directly

Extending for when you need a request-response



Questions?

Absolutely anything: Kafka, gRPC, Terraform...

Can ask me after too.

Where to go next

- Don't want to manage Kafka? Confluent Cloud:
<https://www.confluent.io/confluent-cloud/>
- github.com/travisjeffery
- twitter.com/travisjeffery
- medium.com/@travisjeffery