

Building Confluent Cloud

Travis Jeffery

- twitter.com/travisjeffery
- github.com/travisjeffery
- medium.com/@travisjeffery

Talk roadmap

How to build services on multiple regions, multiple clouds in a manageable way.

- What is Confluent Cloud
- How we built it
- Lessons learned
- Where we're heading
- Ask Me Anything

Travis Jeffery

- Cloud Eng at Confluent
- Wrote my own Kafka in Go called Jocko
- Other projects: Timecop, Mocha, Clang-Format on Xcode
- Worked at Basecamp, Segment

My face on the internet ->



What is Confluent Cloud?

- Today: Kafka as a Service
- Available on AWS and GCP in many regions

The screenshot shows the 'New cluster' configuration page in the Confluent Cloud management console. The page is divided into a left sidebar, a main configuration area, and a right summary panel.

Left Sidebar:

- MANAGEMENT**
 - Clusters
- RESOURCES**
 - Documentation
 - Community
 - Contact us

Main Configuration Area:

- General**
 - Cluster name*: Death Star
 - Write throughput per second: 1MB (slider from Low to High)
 - Read throughput per second: 1MB (slider from Low to High)
 - Total storage: 500GB (slider from Low to High)
 - Cloud provider*: Amazon Web Services (selected) / Google Cloud Platform
 - Region*: us-east-1
 - High Availability (HA)*: Single zone

Right Summary Panel:

- Price:** \$0.6849 PER HOUR
- Billing cycle:** First payment
- SUMMARY:**
 - Price (per hour)
 - Peak (in)
 - Peak (out)
 - Total storage
 - Cloud
 - Region
 - Durability

Bottom Bar:

- Empire Professional** (Travis Jeffery)
- Buttons:** Continue, Cancel

What is Confluent Cloud?

Tomorrow:

- More services: Schema Registry, Connect, Streams, Confluent Control Center
- More features: Topic management, dashboard metrics, SSO auth, etc.
- More regions, more clouds

Building it

The goal: Try to build a PaaS on multiple regions, multiple clouds in a manageable and cost efficient way.

Starting point

- CLI in Python
- Orchestrated by Kubernetes
- Generated YAML with Jinja
- Shelled out to kubectl


Problems with this setup

- No API to build on
- No UI - no user access, staff managed
- No infrastructure management
- No version consensus
- Hard to test
- Static configuration
- Lots of work to onboard to dev or use it

Fresh start

- Libraries/APIs first
- Type-safe calls via Go and gRPC
- Kafka for async, secure messaging to other regions
- Terraform managed infrastructure/configuration
- Still orchestrated by Kubernetes

Request flow



Libraries first, then APIs, then CLIs

- Helps you focus on flexible, robust API
- Can put aside requirements of end program
- End program is a small layer tying together config and libs
- Accessed via HTTP/RPC and CLI

Go and gRPC

- Same lang as our infra, tighter integration and clients: Kubernetes, Terraform, Docker
- Type-safe calls
- Easy to run different API versions
- Defined/managed in protocol buffers
- Service clients for free

Kafka

- Cross region, cross cloud, simple, and secure networking
- Central cluster in mothership
- SASL/PLAIN authentication
- All services just need to know its endpoint on the internet

Routing messages per region

```
MSG = {  
  ROUTER = {  
    CLOUD: 'AWS',  
    REGION: 'US-WEST-2'  
  }  
  DATA: ...  
}
```

→ CREATE_ACCOUNT_TOPIC

```
REGIONAL SERVICE  
MSG ← CONSUMER  
IF MSG IS FOR ME  
  OPERATE  
ELSE  
  IGNORE
```

Terraform

- Provisions infrastructure
- Ties configuration
- Secrets stored in/looked up from KMS

All it takes to add a new region ->

```
module "k8s-sz-a1" {  
    docker_repo = "${var.docker_repo}"  
    azs         = ["ap-southeast-1a"]  
    dd_api_key  = "${var.datadog_api_key}"  
    env         = "${var.caas_env}"  
    cloud       = "aws"  
    region      = "ap-southeast-1"  
    caas_domain = "${var.caas_domain}"  
}
```

Billing

- Using Stripe for payments
- Subscription API was too limited - no postpay billing
- Wrote our own billing service
- Event table stores each change user made on their clusters with associated price per second
- Job runs next month, sums total from events, bills

Bill item example

- Cluster created March 15th at a \$0.10 price per second
 - Cluster updated March 16th to a \$0.20 price per second
 - Cluster deleted March 17th
- => $[0.10 * (24 * 60 * 60)] + [0.20 * (24 * 60 * 60)].$

What's next

- Schema Registry, Connect, Streams, Confluent Control Center
 - Metrics: Traffic metrics for users via Kafka and showno in our UI
- Tooling: Internal services for support, on-call, etc.
- Testing: Integration and UI

**Ask Me
Anything**

Where to go

- Blog posts and open source from what we're building
- github.com/travisjeffery
- medium.com/@travisjeffery
 - medium.com/@travisjeffery
 - github.com/travisjeffery
- Get in touch if you're interested in working on this